

Network setup and training conditions

In this notebook we show how the models were built and trained. All settings except:

- data source (reduced data quality + number of features)
- number of training epochs

are identical with those presented in the paper. Running this notebook will not yield the performance numbers presented in the paper because the beefy networks require larger variation in the data.

The AttA3 network would probably require few days to train so running it with full epoch count is not advisable if one wants a quick look at the models.

In [1]:

```
%load_ext autoreload
%autoreload 2
# %matplotlib widget
import matplotlib
import matplotlib.pyplot as plt
from matplotlib import cm
import numpy as np
import pandas as pd
from pathlib import Path

import torch
from fastai.data.transforms import TfmdLists, DataLoaders
import fastai.learner
import fastai.callback.schedule
import fastai.callback.tracker
import fastai.basics
import fastai.losses

import network_definitions
```

In [2]:

```
data_load_path = Path("../data_sample")
train_data = np.load(data_load_path / "train_samples.npy")
valid_data = np.load(data_load_path / "valid_samples.npy")
train_samples = list(train_data) # to match the requirements of fastai.TfmdLists
valid_samples = list(valid_data)

sample_length = train_data[0].shape[0]
# Settings to match the paper experiments. Except the number of columns, all the
future_length = 60
past_length = sample_length - future_length
no_features = train_data[0].shape[1]
known_column_indexes = list(range(no_features)) # [0, 1, 2, 3]
command_indexes = [2]
target_indexes = [3]
feature_groups = [(0, 1, 3)]
average_range = 0.2
batch_size = 32

feature_itemizer = network_definitions.FeatureItemizer(future_length, known_colu
```

```

                                command_indexes, target_indexes, feature_indexes)
tls_train = TfmdLists(train_samples, [feature_itemizer])
tls_valid = TfmdLists(valid_samples, [feature_itemizer])
data_loader = DataLoaders.from_dsets(tls_train, tls_valid, bs=batch_size, drop_last=True)
print(f>Data volumes: Train: {len(train_samples)}, Validation: {len(valid_samples)}<

```

Data volumes: Train: 7255, Validation: 1073

In [3]:

```

def get_DAffAffGau():
    model = network_definitions.ConstructDelayNet(no_features, len(command_indexes),
                                                    filter_low_classname="AffineTransformer",
                                                    aggregator_low_expansion=1, aggregator_high_expansion=1,
                                                    temporal_contractor_classname="TemporalContractor",
                                                    filter_high_classname="GaussFilter",
                                                    aggregator_high_expansion=1, aggregator_low_expansion=1)

    return model

def get_DLogAffGau():
    model = network_definitions.ConstructDelayNet(no_features, len(command_indexes),
                                                    filter_low_classname="LogGaussFilter",
                                                    aggregator_low_expansion=1, aggregator_high_expansion=1,
                                                    temporal_contractor_classname="TemporalContractor",
                                                    filter_high_classname="GaussFilter",
                                                    aggregator_high_expansion=1, aggregator_low_expansion=1)

    return model

def get_AttA1():
    model = network_definitions.ICCP_Wrap_AttentionModel(no_features, len(command_indexes),
                                                         hidden_size=8, num_layers=1)

    return model

def get_AttA3():
    model = network_definitions.ICCP_Wrap_AttentionModel(no_features, len(command_indexes),
                                                         hidden_size=512, num_layers=1)

    return model

def instantiate_learner(model, data_loader):
    learner = fastai.Learner(data_loader, model, loss_func=fastai.losses.cross_entropy,
                             cbs=[fastai.callback.tracker.ReduceLROnPlateau(patience=20,
                                     fastai.callback.tracker.EarlyStoppingCallback(patience=10),
                                     fastai.callback.tracker.SaveModelCallback(fname="best_model.pkl"),
                                     fastai.callback.tracker.TerminateOnNaNCallback(),
                                     fastai.callback.progress.CSVLogger("learning_progress.csv"),
                                     ],)

    return learner

def evaluate_learner(crt_learner, samples):
    raw_preds, raw_targets = crt_learner.get_preds()
    eval_itemset_arr = np.array(samples)
    feature_itemizer = crt_learner.dls[0].fs[0]
    future_temp_pred_transf = network_definitions.decode_predictions_from_the_network_outputs(raw_preds, feature_itemizer)
    future_temp_target_transf = network_definitions.decode_predictions_from_the_network_outputs(raw_targets, feature_itemizer)
    mae = np.average(np.abs(future_temp_target_transf[:, -future_length:] - future_temp_pred_transf[:, -future_length:]))
    mae_wrt_to_time = np.average(np.abs(future_temp_target_transf[:, -future_length:] - future_temp_pred_transf[:, -future_length:]))
    return mae, mae_wrt_to_time, future_temp_pred_transf, future_temp_target_transf

```

```
In [4]: model_names = []
        learners = []
```

```
In [5]: # Experimental setup in paper
        # global_lr_rate = 1e-3
        # global_no_max_epochs = 1000

        # Experimental setup here, for demonstration
        global_lr_rate = 1e-3
        global_no_max_epochs = 3
```

```
In [6]: DAffAffGau_learner = instantiate_learner(get_DAffAffGau(), data_dloader)
        DAffAffGau_learner.fit_one_cycle(global_no_max_epochs, global_lr_rate)
        model_names.append("DAffAffGau")
        learners.append(DAffAffGau_learner)
```

epoch	train_loss	valid_loss	time
0	0.011518	0.017815	00:12
1	0.008342	0.010210	00:12
2	0.007646	0.009017	00:12

Better model found at epoch 0 with valid_loss value: 0.01781485415995121.
Better model found at epoch 1 with valid_loss value: 0.010210155509412289.
Better model found at epoch 2 with valid_loss value: 0.009016791358590126.

```
In [7]: DLogAffGau_learner = instantiate_learner(get_DLogAffGau(), data_dloader)
        DLogAffGau_learner.fit_one_cycle(global_no_max_epochs, global_lr_rate)
        model_names.append("DLogAffGau")
        learners.append(DLogAffGau_learner)
```

epoch	train_loss	valid_loss	time
0	0.009463	0.088862	00:13
1	0.007932	0.023762	00:14
2	0.007335	0.007523	00:13

Better model found at epoch 0 with valid_loss value: 0.0888623520731926.
Better model found at epoch 1 with valid_loss value: 0.023762168362736702.
Better model found at epoch 2 with valid_loss value: 0.007523125037550926.

```
In [8]: AttA1_learner = instantiate_learner(get_AttA1(), data_dloader)
        AttA1_learner.fit_one_cycle(global_no_max_epochs, global_lr_rate)
        model_names.append("AttA1")
        learners.append(AttA1_learner)
```

epoch	train_loss	valid_loss	time
0	0.031977	0.009789	00:47
1	0.009757	0.011009	00:45
2	0.008705	0.010445	00:47

Better model found at epoch 0 with valid_loss value: 0.00978886429220438.

```
In [9]: AttA3_learner = instantiate_learner(get_AttA3(), data_dloader)
AttA3_learner.fit_one_cycle(global_no_max_epochs, global_lr_rate)
model_names.append("AttA3")
learners.append(AttA3_learner)
```

epoch	train_loss	valid_loss	time
0	0.008202	0.008381	01:12
1	0.006932	0.006418	01:09
2	0.006595	0.006736	01:09

Better model found at epoch 0 with valid_loss value: 0.008381090126931667.
 Better model found at epoch 1 with valid_loss value: 0.006417920347303152.

Network summary.

Note where the bulk of parameters is concentrated in each network.

```
In [10]: for k in range(len(learners)):
          print(f"Architecture for {model_names[k]}:")
          network_definitions.print_model_weights_rec(learners[k].model, max_level=1)
```

Architecture for DAffAffGau:

```
<class 'network_definitions.ConstructDelayNet'> with 12049 parameters:
  <class 'network_definitions.BankedFilters'> with 64 parameters:
    Parameters: 64
  -----
  <class 'network_definitions.FeatureAggregationStack'> with 664 parameters:
    Parameters: 664
  -----
  <class 'network_definitions.AffineTransform'> with 32 parameters:
    Parameters: 32
  -----
  <class 'network_definitions.BankedFilters'> with 776 parameters:
    Parameters: 776
  -----
  <class 'network_definitions.FeatureAggregationStack'> with 10513 parameter
s:
  Parameters: 10513
  -----
Parameters: 12049
-----
```

Architecture for DLogAffGau:

```
<class 'network_definitions.ConstructDelayNet'> with 13681 parameters:
  <class 'network_definitions.BankedFilters'> with 64 parameters:
    Parameters: 64
  -----
  <class 'network_definitions.FeatureAggregationStack'> with 2296 parameter
s:
  Parameters: 2296
  -----
  <class 'network_definitions.AffineTransform'> with 32 parameters:
    Parameters: 32
  -----
  <class 'network_definitions.BankedFilters'> with 776 parameters:
    Parameters: 776
  -----
  <class 'network_definitions.FeatureAggregationStack'> with 10513 parameter
s:
  Parameters: 10513
  -----
```

Parameters: 13681

Architecture for AttA1:

<class 'network_definitions.ICCP_Wrap_AttentionModel'> with 8105 parameters:

<class 'network_definitions.AttentionModel'> with 8105 parameters:

Parameters: 8105

Parameters: 8105

Architecture for AttA3:

<class 'network_definitions.ICCP_Wrap_AttentionModel'> with 16239929 parameters:

<class 'network_definitions.AttentionModel'> with 16239929 parameters:

Parameters: 16239929

Parameters: 16239929

Performance evaluation

There is no expectation that these numbers will match the paper's results. The networks are identical but the data quality is lower.

In [11]:

```
maes = []
maes_timewise = []
predictions = []

for l in learners:
    mae, mae_t, scaled_preds, scaled_targets = evaluate_learner(l, valid_samples)
    maes.append(mae)
    maes_timewise.append(mae_t)
    predictions.append(scaled_preds)

# Adding "zero" predictor
raw_preds, raw_targets = learners[0].get_preds()
zero_preds = torch.zeros_like(raw_preds)
eval_itemset_arr = np.array(valid_samples)
feature_itemizer = learners[0].dls[0].fs[0]
zero_transf = network_definitions.decode_predictions_from_the_network(zero_preds)
mae_zero = np.average(np.abs(scaled_targets[:, -future_length:] - zero_transf[:, -future_length:]))
mae_zero_time = np.average(np.abs(scaled_targets[:, -future_length:] - zero_transf[:, -future_length:]))
model_names.append("Zero")
maes.append(mae_zero)
maes_timewise.append(mae_zero_time)
predictions.append(zero_transf)
model_names = model_names[:len(learners)+1] # just to make sure that on re-runs
for k in range(len(model_names)):
    print(f"Performance of {model_names[k]}: MAE: {maes[k]:.4}")
```

Performance of DAffAffGau: MAE: 0.1356

Performance of DLogAffGau: MAE: 0.1133

Performance of AttA1: MAE: 0.1457

Performance of AttA3: MAE: 0.09819

Performance of Zero: MAE: 0.1366